

Project 1 - Explainable AI in Radiology

Utrecht University

Arda Adali 0385433 Lena Feiler 1751611 Kerem Dogan 1696637 Sofia Bianchini 7273746

Abstract

This study applies explainable AI techniques to detect pleural effusion in chest X-rays using a DenseNet169 model. Saliency Maps highlight image regions relevant to predictions, while Example-Based Explanations retrieve similar cases for comparison. Explanation quality is assessed using Quantus for Saliency Maps and Normalized Discounted Cumulative Gain (nDCG) for case rankings. The explanations presented in this work show good initial results but would still require improvement before being applied in the field.

Keywords: Chest X-Ray, Pleural Effusion, Explainable AI, Saliency maps, GradCAM, Example-based explanations, Cosine Similarity, nDCG, DenseNet

1. Introduction

Machine learning is increasingly used to solve complex problems across domains. In healthcare AI supports diagnostic and treatment decisions (9). However, many models act as black boxes with limited insight into how they generate outputs (2), making explainable AI essential to improve transparency and trust. X-ray analysis is one of the main applications of AI in medicine, especially for detecting bone and lung diseases (14). This paper presents a machine learning model to predict pleural effusion from CXR images.

2. Methods & Materials

Dataset and Preprocessing

The dataset used for this study was a smaller version of the CheXpert dataset (5), containing frontal and lateral chest X-rays, of which only the frontal images were used. Among the pathology labels available, particular attention was paid to *Pleural Effusion*, which served as the classification target.

Images with missing or uncertain *Pleural Effusion* labels were removed to improve reliability and all images were resized to 224×224 pixels and normalized to the $[0, 1]$ range. To address the class imbalance, class weights were calculated based on label frequencies and applied during training, where an 80-20 train test split was used. Since the original X-rays were grayscale, each image was duplicated across three channels to match the input format required by ImageNet-pretrained models.

Model and Training

A DenseNet169 (12) convolutional neural network was used for the *Pleural Effusion* classification. The model was initialized with ImageNet (8) weights, freezing the lower layers while fine-tuning upper layers. A global average pooling layer was added, followed by dense layers with ReLU activations. The final layer used a sigmoid activation for binary classification. Training was performed using the Adam optimizer with a learning rate of 0.0001, binary cross-entropy loss, and class weighting. The model was trained for 10 epochs and saved.

Explanations

For this experiment two different explanations were exploited. This report explores Gradient-weighted Class Activation Mapping (**Grad-CAM**), which identifies critical regions in an image by leveraging gradients in the last convolutional layer of a CNN and global-average-pooling to assess neuron importance (15). Grad-CAM generates heatmaps that highlight relevant features for model predictions. To evaluate the quality of saliency maps, the Quantus toolkit (11) was used, employing metrics such as Pixel-Flipping

(6) (Faithfulness), Local Lipschitz Estimate (3) (Robustness), Model Parameter Randomisation (1) (Randomisation), Relevance Mass Accuracy (4) (Localisation), and Complexity (7) (Complexity).

For example-based explanations, image feature embeddings from the last pooling layer are extracted and ranked using **cosine similarity**, measuring vector orientation and magnitude (10). These rankings are compared to expert-defined rankings via Normalized Discounted Cumulative Gain (**nDCG**) (13), which assesses ranking quality by assigning relevance scores in the range [5.5, 1], computing discounted gain, and normalizing against the ideal expert ranking. This ensures that the model groups similar images in a clinically meaningful way.

3. Results & Discussion

The following sections present and discuss the results of testing the model and explanation methods from Section 2 on predefined test images.

Model Training and Evaluation




Figure 1: Training and validation per epoch

During training, the DenseNet169 model steadily improved in training accuracy and loss, though validation accuracy fluctuated, indicating some overfitting (Figure 1). The final model achieved 79.21% accuracy on the test set, with a precision of 63.41% and recall of 81.25%, showing good detection of positive *Pleural Effusion* cases. The F1-score was 71.23%, balancing precision and recall, while the ROC AUC reached 84.96%, demonstrating strong overall discriminatory power. A binomial test yielded a p-value of 2.03×10^{-17} , confirming the performance is significantly better than chance.

Saliency Maps

Figure 3 shows the GradCAM explanations generated to gain insights about the models' focus on the test images. The most important region, as identified by the model, is represented in seismic colours: red, white, and blue. The bounding boxes on the explanation boxes circle the critical parts of the image, that the explanation should focus on. In all five test images, the explanations highlight parts of the bounding




Figure 2: GradCAM Explanations

box but also often highlight irrelevant areas outside of it.

In most cases, the heatmaps focus on the lower lung regions associated with Pleural Effusion, suggesting the model identifies clinically relevant areas. However, some explanations highlight unrelated regions, indicating that the model occasionally relies on uninformative features to classify Pleural Effusion.

Additionally, Table 3 shows the mean and standard deviation of the Quantus evaluation metrics on the test images ¹.

	Mean	Standard deviation
Complexity	9.360562	0.241967
Local Lipschitz Estimate	0.032051	0.031284
Pixel Flipping	1	0

Table 1: Quantus: Explanation Evaluation

The complexity metric (9.3) reflects the expected high complexity of CNN explanations. The robustness score (0.032051) indicates the model is highly stable to small input perturbations. Finally, the faithfulness measure confirms that removing important pixels causes a significant drop in confidence, showing the explanation accurately highlights the model’s key features.

Example-based Explanations

The results show strong alignment between the scientists’ and cosine similarity rankings. The high nDCG score (0.96) confirms that extracted features correlate well with expert judgment. While the line plot (Figure 3) reveals some ranking mismatches, overall similarity remains consistent, especially in the top-ranked images. The heatmap in plot (Figure 4) further shows that five of the top six images appear in both rankings, though in a slightly different order.




Figure 3: The line plot of images ranked by the scientists and cosine similarity.




Figure 4: heatmap of the similarity rankings (identical images share the same color)

4. Conclusion

In the medical field, understanding the decision process of machine learning models is essential for transparency and trust in the findings. Thus, increasing the importance of integrating explanations in assistance systems provided to medical workers. In this study, two explanation methods were presented for this purpose and assessed for their quality. Overall, the results showed promising outcomes, especially in terms of having a close ranking compared to a professional. However, improvements upon the model and explanation could be made with additional time and different classification models should be compared in the future, preferably using the PyTorch environment for better integration with Quantus.

1. Please note that the other two metrics (Randomisation and Localisation) were left out of the table due to practical errors between the TensorFlow model and the Quantus explanations, which is explained in Appendix A.

Appendix A. Quantus Error Messages

While Quantus was adapted to TensorFlow, it appears that two of the requested methods are not fully compatible. We have looked for help but could not fix it until the deadline and had not trained a new model since we were initially convinced, that with the promise of compatibility, the method should work, at least with the in-built explanation methods of quantus - which it did not. In Figure 5 one can see the error displayed when applying the localisation metric (Relevance Mass Accuracy) to the in-built explanation method of Integrated Gradients. The “inputs should share same dimensions” error presents itself even when the shapes of the two specified inputs are clearly the same.

Figure 5: The error displayed when for the localisation metric on the in-built Integrated Gradients explanation method with a TensorFlow model.

```

✓ [50] print(x_batch.shape == s_batch.shape)
0s
➔ True

[51] # 5. Localisation: relevance mass accuracy (Arras et al., 2021)
0s
localisation = quantus.RelevanceMassAccuracy(
)(model=model,
  x_batch=x_batch,
  y_batch=y_predicted,
  a_batch=None,

  #explain_func= custom_explain,
  explain_func=quantus.explain,
  explain_func_kwargs={"method": "IntegratedGradients"},
  s_batch=s_batch,
)

print(f"Localisation Scores (relevance mass accuracy): {localisation}")

-----
AssertionError                                Traceback (most recent call last)
<ipython-input-51-74a0a3fbf4b3> in <cell line: 0>()
      1 # 5. Localisation: relevance mass accuracy (Arras et al., 2021)
----> 2 localisation = quantus.RelevanceMassAccuracy(
      3 )(model=model,
      4     x_batch=x_batch,
      5     y_batch=y_predicted,

-----
4 frames -----
/usr/local/lib/python3.11/dist-packages/quantus/helpers/asserts.py in assert_segmentations(x_batch, s_batch)
    219 ), "The inputs 'x_batch' and segmentations 's_batch' should include the same number of samples."
    220 assert (
--> 221     np.shape(x_batch)[2:] == np.shape(s_batch)[2:]
    222 ), "The inputs 'x_batch' and segmentations 's_batch' should share the same dimensions."
    223 assert (

AssertionError: The inputs 'x_batch' and segmentations 's_batch' should share the same dimensions.

```

References

- [1] Julius Adebayo, Justin Gilmer, Michael Muelly, Ian Goodfellow, Moritz Hardt, and Been Kim. Sanity checks for saliency maps. *Advances in neural information processing systems*, 31, 2018.
- [2] Saad Bin Ahmed, Roberto Solis-Oba, and Lucian Ilie. Explainable-ai in automated medical report generation using chest x-ray images. *Applied Sciences*, 12(22):11750, 2022.
- [3] David Alvarez Melis and Tommi Jaakkola. Towards robust interpretability with self-explaining neural networks. *Advances in neural information processing systems*, 31, 2018.

- [4] Leila Arras, Ahmed Osman, and Wojciech Samek. Clevr-xai: A benchmark dataset for the ground truth evaluation of neural network explanations. *Information Fusion*, 81:14–40, May 2022. ISSN 1566-2535. doi: 10.1016/j.inffus.2021.11.008. URL <http://dx.doi.org/10.1016/j.inffus.2021.11.008>.
- [5] Ashery. Chexpert dataset (downsized version), 2021. URL <https://www.kaggle.com/datasets/ashery/chexpert/data>. Accessed: 2025-03-07.
- [6] Sebastian Bach, Alexander Binder, Grégoire Montavon, Frederick Klauschen, Klaus-Robert Müller, and Wojciech Samek. On pixel-wise explanations for non-linear classifier decisions by layer-wise relevance propagation. *PloS one*, 10(7):e0130140, 2015.
- [7] Umang Bhatt, Adrian Weller, and José MF Moura. Evaluating and aggregating feature-based model explanations. *arXiv preprint arXiv:2005.00631*, 2020.
- [8] Jia Deng, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. Imagenet: A large-scale hierarchical image database. In *2009 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 248–255. IEEE, 2009.
- [9] Izegbua E. Ihongbe, Shereen Fouad, Taha F. Mahmoud, Arvind Rajasekaran, and Bahadar Bhatia. Evaluating explainable artificial intelligence (xai) techniques in chest radiology imaging through a human-centered lens. *Plos one*, 19(10):e0308758, 2024.
- [10] Xinting Han, Xiang Wang, Chong Huang, Chen Xu, and Xiangnan Ye. A survey on metric learning and its applications. *Neurocomputing*, 472:55–73, 2022.
- [11] Anna Hedström, Leander Weber, Daniel Krakowczyk, Dilyara Bareeva, Franz Motzkus, Wojciech Samek, Sebastian Lapuschkin, and Marina M-C Höhne. Quantus: An explainable ai toolkit for responsible evaluation of neural network explanations and beyond. *Journal of Machine Learning Research*, 24(34):1–11, 2023.
- [12] Gao Huang, Zhuang Liu, Laurens Van Der Maaten, and Kilian Q. Weinberger. Densely connected convolutional networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4700–4708, 2017. doi: 10.1109/CVPR.2017.243. URL <https://doi.org/10.1109/CVPR.2017.243>.
- [13] Kalervo Järvelin and Jaana Kekäläinen. Cumulated gain-based evaluation of ir techniques. *ACM Transactions on Information Systems (TOIS)*, 20(4):422–446, 2002.
- [14] Miquel Miró-Nicolau, Gabriel Moyà-Alcover, and Antoni Jaume-i Capó. Evaluating explainable artificial intelligence for x-ray image analysis. *Applied Sciences*, 12(9):4459, 2022.
- [15] Ramprasaath R Selvaraju, Michael Cogswell, Abhishek Das, Ramakrishna Vedantam, Devi Parikh, and Dhruv Batra. Grad-cam: Visual explanations from deep networks via gradient-based localization. In *Proceedings of the IEEE international conference on computer vision*, pages 618–626, 2017.